

DAVIS: A **Depth-Only** End-to-End **Active-Vision** Framework for **HumanoId Soccer Skills**

Anonymous Submission




Fig. 1. **Overview of our humanoid soccer system.** (a) The robot performs goal-oriented shooting by approaching the ball and kicking it into the goal. (b) The robot demonstrates agile multi-directional dribbling with continuous ball control, including turning and weaving around poles. (c) Our active vision design enables the robot to actively reorient its head in yaw and pitch to keep the ball near the center of the field of view, thereby improving perception-action coordination.

Abstract—Humanoid soccer contact skills require more than producing high-impact foot-ball contacts: the robot must close the loop over perception, approach, alignment, impact, and recovery while its own motion induces substantial viewpoint changes, frequent loss of the ball from view, and uncertain contact outcomes. In this work, we ask a compact yet stricter question: can a humanoid learn soccer contact skills using only a head-mounted depth image, proprioceptive history, and an optional low-dimensional task command, and directly output 25-DoF joint PD targets without extra runtime perception or planning modules? To this end, we propose DAVIS, a depth-only end-to-end framework for humanoid soccer skills that learns visibility-aware auxiliary geometry during training, and combines GT-to-prediction annealing, task curricula, and AMP-style motion priors to smoothly bridge privileged supervision and real deployment. Built on this framework, we instantiate representative soccer contact skills, including goal-directed shooting and directional dribbling, through task-specific definitions of objects, commands, rewards, and curricula, and validate them through simulation, Noetix E1 real-robot experiments, and ablations. Supplementary material, additional details and result videos are available at <https://motu1734.github.io/DAVIS/>.

I. INTRODUCTION

Humanoid robots have made substantial progress in reinforcement-learning-driven dynamic locomotion [1]–[6], but soccer remains a harder setting because it requires the robot to perceive, approach, align, contact, and recover while the **target and the robot itself are moving**. Unlike standard locomotion benchmarks, soccer contact skills require the policy to reason not only about *where the ball is*, but also about *how to approach it*, *how to keep the target in view*, and *how contact should change the subsequent ball motion*.

On real humanoids, onboard sensing and whole-body control further amplify these difficulties: the head-mounted view changes with body motion, active head yaw and pitch determine what evidence the policy receives, and small perception or stance errors can change the contact outcome [7]–[11]. Existing soccer systems often manage this complexity with modular perception, localization, planning, and skill-switching interfaces, including RGB detectors, LiDAR/RGB-D-based localization, or external state inputs. **Because soccer contact skills depend on local ball and goal geometry, head-**

mounted depth provides a direct geometric observation from which to study control without externally supplied scene state. This motivates the question: **can a humanoid learn deployable soccer contact skills from only head-mounted depth, proprioceptive history, and low-dimensional task commands — without any external scene state?**

We propose **DAVIS**, a depth-only end-to-end reinforcement learning framework to answer this question. At deployment, the policy remains a closed-loop mapping from head-mounted depth, proprioceptive history, and low-dimensional commands to 25-DoF joint PD targets. This is achieved by placing structure inside the learning signal rather than in the deployment pipeline, following the idea of using privileged simulation information to guide deployable policies [12]–[14]. Auxiliary geometry and visibility predictions are learned from privileged simulation supervision, while the actor gradually transitions from stable geometric guidance to self-predicted features via GT-to-prediction annealing and visibility gating. With this interface fixed, different soccer behaviors can be learned by changing the local objective, geometry, and supervision that define the task.

DAVIS is designed to support a range of soccer skills through a shared learning framework: new skills are specified by task objects, commands, rewards, and curricula while retaining the depth-to-control interface. We evaluate goal-directed shooting and directional dribbling in detail as representative instances of one-shot striking and sustained ball control, using them to assess the feasibility of the framework across complementary contact regimes. Additional examples, including stopping the ball, obstacle-aware dribbling, and ball-loss recovery, are shown in Fig. 5 and the supplementary video. Together, these instances provide evidence of the framework’s extensibility through task-specific training.

In summary, our contributions are as follows: (1) We introduce a **depth-only end-to-end RL framework** for humanoid soccer contact skills, where the deployed policy maps head-mounted depth, proprioceptive history, and low-dimensional task commands directly to 25-DoF joint targets with active-vision head control. (2) We develop a **visibility-aware auxiliary geometry mechanism** that learns task-relevant object geometry from the depth latent, and combines visibility masking with GT-to-prediction annealing to stabilize visual learning without adding runtime perception modules. (3) We instantiate multiple soccer skills within the same framework and establish a complete pipeline from IsaacLab training to E1 real-robot deployment, including depth alignment, policy export, and 25-DoF PD control. Real-robot evaluations and additional skill demonstrations support the framework’s practical feasibility and applicability across soccer tasks.

II. RELATED WORK

A. Robot Soccer Skill Learning

Robot soccer has long served as a challenging benchmark for robot control, requiring the integration of visual perception, real-time decision making, agile locomotion, and

ball-oriented skills [15]–[17]. RoboCup-style soccer systems often relied on modular perception-control pipelines, where sensing, state estimation, and soccer control were separately engineered rather than learned end-to-end [18]. Recent learning-based methods have enabled legged robots to acquire soccer skills through various paradigms, including reinforcement learning and imitation learning [19]–[26]. Yet most existing systems still depend on RGB-based object detection or multi-sensor perception pipelines [27]–[30].

In contrast, our work studies depth-driven humanoid soccer skill learning, where the robot learns to chase, dribble, and kick using only depth observations and proprioception.

B. Whole-Body Loco-Manipulation

Whole-body loco-manipulation on humanoid robots is inherently challenging, as the robot must simultaneously coordinate balance, locomotion, and task-directed manipulation within a high-dimensional and tightly coupled control space [31]–[33]. To make the problem tractable, prior systems often introduce decoupled or hierarchical interfaces, such as separate upper- and lower-body controllers or high-level task policies paired with low-level whole-body controllers [34]–[38]. Humanoid soccer further amplifies this difficulty: unlike manipulating static objects, the ball is a highly dynamic target, requiring the robot to approach from a suitable direction, maintain visibility, and generate precise foot–ball contacts for effective ball control.

To address these challenges, we propose a whole-body control framework, **DAVIS**, that jointly optimizes active visual tracking, approach, and foot–ball interaction for dynamic ball-centric loco-manipulation.

III. METHOD

A. Problem Formulation

We model humanoid soccer contact skills as a POMDP problem [39], denoting the process as $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \gamma \rangle$. The simulator maintains a full state $\mathbf{s}_t \in \mathcal{S}$ containing robot, ball, goal, and contact-related physical variables; during training, privileged variables derived from this state, such as ball pose, goal pose, visibility, contact state, and reference motion, are used as critic inputs, reward-computation terms, and auxiliary-supervision labels. The deployed policy instead observes only a head-mounted depth image $\mathbf{D}_t$, a proprioceptive history $\mathbf{p}_{t-H:t}$ of length H , and an optional low-dimensional task variable $\mathbf{c}_t$; for instances without an external command, $\mathbf{c}_t$ is omitted. We write the deployed observation and action as $\mathbf{o}_t = (\mathbf{D}_t, \mathbf{p}_{t-H:t}, \mathbf{c}_t)$ and $\mathbf{a}_t = \pi_\theta(\mathbf{o}_t) \in \mathbb{R}^{25}$. The proprioceptive history is encoded inside the policy to infer latent body dynamics, while all explicit task geometry used by the actor must be produced from deployable inputs. The low-level controller converts the action into joint PD targets $\mathbf{q}_t^{\text{des}} = \mathbf{q}^0 + \mathbf{s} \odot \mathbf{a}_t$, where $\mathbf{q}^0$ is the nominal joint configuration and $\mathbf{s}$ is the per-joint action scale. The E1 humanoid consists of 23 body DoFs and 2 head DoFs, so we decompose the action as $\mathbf{a}_t = [\mathbf{a}_t^{\text{body}}, \mathbf{a}_t^{\text{head}}]$ with $\mathbf{a}_t^{\text{body}} \in \mathbb{R}^{23}$ and $\mathbf{a}_t^{\text{head}} \in \mathbb{R}^2$. Because the depth camera is mounted on the actuated head, the head action




Fig. 2. **Overview of our DAVIS framework.** The observation stream contains the command c_t , a **single** 168×80 **depth image**, proprioceptive history, and privileged simulation states. The structured-depth module encodes the depth input into a 32-D latent, predicts ball/goal geometry and visibility confidence through **auxiliary heads**, forms confidence-gated geometry features, and combines them with a 64-D HIM/history feature. The actor receives the command, depth latent, gated geometry/confidence features, proprioceptive history, and HIM feature, and outputs a 25-D position residual for the PD controller. The auxiliary losses, GT-to-prediction annealing, PPO update, task reward, and style reward use simulation-only signals from IsaacLab rollouts, including ball, goal, and visibility ground truth. Dashed boxes in the structured-depth module indicate optional task-specific modules.

is part of the sensing loop: it directly changes the next depth observation. We optimize the asymmetric actor-critic policy with PPO [40] by maximizing the expected discounted return

$$\max_{\theta} \mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^T \gamma^t (r_t^{\text{task}} + \lambda_{\text{amp}} r_t^{\text{amp}} - p_t^{\text{reg}}) \right], \quad (1)$$

where r_t^{task} denotes task reward, r_t^{amp} encourages motion near a reference prior, and p_t^{reg} penalizes joint limits, action changes, unstable posture, collisions, and unsafe contact.

B. Depth-Only End-to-End Structured Soccer RL Framework

We build **DAVIS** around a deployment interface that remains deliberately narrow: the actor receives head-mounted depth, proprioceptive history, and an optional low-dimensional command, and outputs 25-DoF joint target residuals for the PD controller, as illustrated in Fig. 2. During training, privileged simulation states provide dense supervision for perception, value learning, and motion regularization; at deployment, the policy keeps only the depth/proprioception-to-control path.

To keep this deployment path compact, we introduce task structure only through training-time supervision and regularization. The depth latent feeds auxiliary geometry and visibility heads, whose predictions are confidence-gated before entering the actor. Simulation ground truth supplies geometry and visibility labels, but only as auxiliary supervision and as early-stage geometric guidance. A GT-to-prediction annealing schedule gradually replaces privileged features with predicted features, so the actor first learns how local geometry should affect whole-body motion and

then operates with geometry inferred from depth. In parallel, PPO optimizes task return, the history estimator learns proprioceptive state inference, and AMP-style motion priors regularize full-body motion. These components make long-horizon soccer contact learning tractable while preserving a single deployable policy graph [41]–[43].

This framework separates reusable learning machinery from task definition. The reusable part consists of the depth/proprioceptive encoders, auxiliary geometry and visibility learning, Asymmetric actor-critic optimization, motion-prior regularization, and depth alignment between simulation and the real camera. A new soccer skill is specified by choosing its task objects, task variables, rewards, curricula, and optional contact conditioning. Once these elements define what the robot should accomplish, the same training pipeline can optimize a policy under the same deployment constraints.

C. Visibility-Aware Auxiliary Geometry

Learning task geometry from depth only through task returns is inefficient: task objects can occupy a small image region, while the gradient from PPO returns must pass through the actor and the depth encoder before shaping visual features. We therefore attach a visibility-aware auxiliary geometry module to the depth latent and supervise it with privileged simulation labels during training. **As shown in Fig. 3, the module predicts object-level geometry and visibility inside the policy, so the actor receives direct object-level supervision without adding a runtime perception module.**

Let $\mathbf{z}_t^D = f_D(\mathbf{D}_t) \in \mathbb{R}^{32}$ be the depth latent. For each




Fig. 3. Scalable auxiliary geometry. Each task object is assigned a geometry head and a visibility head supervised by privileged labels during training; their outputs are visibility-gated and concatenated into the actor feature, enabling the same mechanism to extend from one object to n objects.

task-relevant object $\text{obj}_i \in \mathcal{O}$, the auxiliary module predicts its 3-D center and visibility confidence from the depth latent,

$$(\hat{\mathbf{y}}_t^i, \hat{q}_t^i) = h_{\text{aux}}^i(\mathbf{z}_t^D), \quad \hat{p}_t^i = \sigma(\hat{q}_t^i), \quad (2)$$

where $\mathbf{y}_t^i = [x_t^i, y_t^i, z_t^i]^\top$ denotes the object center in a right-handed robot-local frame, with x along the head-yaw forward direction, y to the left, and z upward. Both geometry and visibility predictions depend only on the current 32-D depth latent; head states are provided separately to the actor through proprioception. The same construction extends to $\mathcal{O} = \{\text{obj}_i\}_{i=1}^n$ by assigning each object an indexed 3-D center prediction $\hat{\mathbf{y}}_t^i$, visibility prediction $\hat{q}_t^i$, ground-truth center $\mathbf{y}_t^i$, and visibility label v_t^i .

Visibility is learned jointly with geometry because a head-mounted camera can temporarily lose task-relevant objects. Geometry is supervised only when the corresponding object is visible, while visibility is trained with a binary objective,

$$\mathcal{L}_{\text{aux}} = \sum_{i=1}^n \lambda_i v_t^i \ell_i(\hat{\mathbf{y}}_t^i, \mathbf{y}_t^i) + \sum_{i=1}^n \lambda_{v_i} \text{BCE}(\hat{q}_t^i, v_t^i). \quad (3)$$

Visibility labels are computed using the full camera transform together with valid field-of-view and depth-range constraints; for multi-point objects, visibility is positive if any predefined structural point is valid, while the geometry target always remains the single 3-D object center.

To avoid forcing the actor to rely on unreliable predictions early in training, we gradually replace privileged geometry with predicted geometry. Let $\beta_k \in [0, 1]$ denote the prediction weight at policy update k . For object i ,

$$\mathbf{y}_{t,k}^{e,i} = (1 - \beta_k) \mathbf{y}_t^i + \beta_k \hat{\mathbf{y}}_t^i. \quad (4)$$

Using the correspondingly annealed visibility confidence $p_{t,k}^{e,i}$, we define the detached confidence gate

$$c_{t,k}^i = \text{clip}(\text{stopgrad}(p_{t,k}^{e,i}), 0.05, 1). \quad (5)$$

The actor-side auxiliary representation is then

$$\mathbf{u}_{t,k}^{\text{aux}} = \text{Concat}_{i=1}^n [c_{t,k}^i \mathbf{y}_{t,k}^{e,i}, c_{t,k}^i]. \quad (6)$$

As β_k increases from 0 to 1, actor transitions from privileged geometric guidance to depth-based predictions. At deployment, all auxiliary features are predicted solely from depth.

D. Skill Instantiation

Once the deployment interface in Sec. III-B is fixed, a soccer skill is specified by a task definition rather than by switching runtime modules. This separation makes the interface reusable: new skills retain the depth-to-control path and are introduced by changing task objects, commands, objectives, and curricula. We denote such a definition by $\mathcal{T} = (\mathcal{O}, \mathbf{c}_t, \mathbf{y}_t, r_t, \mathcal{K})$, where $\mathcal{O}$ is the set of task objects, $\mathbf{c}_t$ is the optional command, $\mathbf{y}_t$ is the task geometry supervised through Sec. III-C, r_t defines the contact objective, and $\mathcal{K}$ is the curriculum. We use goal-directed shooting and directional dribbling as representative instances to verify the framework across complementary contact regimes; the following details illustrate how the shared framework is instantiated for these examples.

1) *Goal-Directed Shooting*: Shooting is defined by the ball-goal geometry. Let $\mathbf{b}_t$ and $\mathbf{g}_t$ denote the ball position and goal center in the robot frame, and let $\mathbf{d}_t^{bg} = \frac{\mathbf{g}_t - \mathbf{b}_t}{\|\mathbf{g}_t - \mathbf{b}_t\|_2 + \epsilon}$ be the normalized ball-to-goal direction. A behind-ball target can then be written as $\mathbf{x}_t^{\text{behind}} = \mathbf{b}_t - \rho \mathbf{d}_t^{bg}$, which turns chasing, alignment, and contact into one geometric target family. In this task, the object set is $\mathcal{O}_{\text{shoot}} = \{\text{ball}, \text{goal}\}$, the task geometry is $\mathbf{y}_t^{\text{shoot}} = (\mathbf{b}_t, \mathbf{g}_t)$, and the objective is to strike the ball so that its exit motion moves toward the goal. The reward is organized around run-up, contact, terminal success, visual maintenance, and safety regularization, while the curriculum gradually broadens the initial ball placement and approach difficulty. Active vision is maintained through the head DoFs: the head keeps the ball within view, and the body gradually follows the viewing direction.

2) *Directional Dribbling*: Dribbling is defined by a direction command and sustained ball regulation. Let c_t be the commanded direction angle, selected from 12 directions uniformly spaced over 360° , and $\mathbf{d}_t^{\text{cmd}} = [\cos c_t, \sin c_t]$ the corresponding horizontal direction. The task object set is $\mathcal{O}_{\text{dribble}} = \{\text{ball}\}$, and the auxiliary geometry is the 3-D ball center $\mathbf{y}_t^{\text{ball}} = \mathbf{b}_t$. The commanded direction $\mathbf{d}_t^{\text{cmd}}$ and the robot-ball horizontal distance d_t are task variables used to define the dribbling objective rather than outputs of the auxiliary geometry head. Unlike shooting, the objective is not a single high-speed strike; the policy must keep the ball within a controllable band and generate progress along $\mathbf{d}_t^{\text{cmd}}$. We express the controllable-band condition by $\chi_t^{\text{band}} = \mathbb{I}(|d_t - d_0| < \Delta_d, v_{lo} < \|\mathbf{v}_t^b\| < v_{hi}, \mathbf{v}_t^b \cdot \mathbf{d}_t^{\text{cmd}} > 0)$. The reward covers commanded progress, re-approach, heading, visual maintenance, and out-of-band penalties. Continuous dribbling further uses contact-intent-conditioned AMP and phase masks so approach, touch, follow-through, and recovery emphasize different motion states.

E. Training and Deployment Pipeline

The deployment policy uses a single-frame head-mounted depth image, a five-step proprioceptive history, and task




Fig. 4. Depth alignment transforms the raw depth frame into the aligned policy input through distortion correction, cropping from 168×94 to 168×80 , ball ROI occlusion, sensor noise, random masks and normalization, so the same single-frame depth interface can be used both in sim and real.




Fig. 5. **Qualitative examples of learned soccer skills.** In simulation: (a) curved and penalty kicks, (b) cone-avoidance dribbling, (d) dribbling across a soccer field, and (f) stopping a rolling ball, where red and blue arrows indicate robot and ball motion, respectively, and the purple box marks the stopped ball. On the real robot: (c) foot-ball contact during dribbling, with arrows indicating ball motion, and (e) a close-up dribbling sequence.

variables to output 25-DoF joint PD targets. Since this depth image is the main sim-to-real interface, we align simulation and deployment before it enters the actor. As shown in Fig. 4, a raw frame D_t is processed to $\tilde{D}_t$. During training, the aligned frame is further perturbed to reduce dependence on clean depth and serve as domain randomization for sim-to-real transfer [44], [45]. Both simulated and real depth frames are cropped to 168×80 , clipped to 0.3–5.0 m, normalized, and zero-filled at invalid pixels. Policies are trained in Isaac Lab for 20,000 PPO iterations with AMP and a history-informed latent estimator (HIM) [46], then exported to ONNX. On the robot, the actor runs at 50 Hz, the low-level PD controller at 200 Hz, and the depth stream at 15 Hz; no external localization, object detector, state estimator, or runtime soccer planner is used.

IV. EXPERIMENTS

We evaluate whether **DAVIS** supports both simulated and real-world soccer contact skills, as illustrated in Fig. 5. And we answer the following questions: **Q1**. Can a depth-only

end-to-end policy perform soccer contact skills? **Q2**. Which training structures are necessary? **Q3**. Does the learned policy transfer to the real Noetix E1 robot?

A. Simulation Experiments

To answer Q1, we extensively evaluate **DAVIS** in IsaacLab simulator on two representative skill categories in soccer: shooting and dribbling.

Shooting and directional dribbling are trained as separate task-specific checkpoints that share the same depth-to-control interface; they are not a single multi-skill policy. At deployment, the shooting actor receives no operator command, while the dribbling actor receives only a joystick direction.

Shooting: The robot is required to perform shooting with randomly initialized robot or ball positions. In the penalty kick task, the robot starts within angle of $[-100^\circ, 100^\circ]$ behind the ball and distances of 2–3.5 m. In the free kick task, the robot is initialized 5 m in front of the goal, and is required to kick the ball within the region $x \in [0, 4]$ m and $y \in [-3, +3]$ m. The difficulty level categories are shown in




Fig. 6. (a) Free-kick and Penalty-kick performance in simulation. Each region uses over 50 trials. Darker colors indicate higher success rates, and dashed lines mark difficulty regions. (b) Footstep trajectories during shooting. The policy shows human-like gait and speed pattern.




Fig. 7. Performance of repeated S-slalom in simulation. (a) Statistical results of 100 S-slalom trajectories. (b) Foot-contact record from a representative successful trajectory.

Goal-Directed Shooting						Directional Dribbling										
Setting	Metric	Easy	Med.	Hard	Total	Setting	Metric	N = 3			N = 5			N = 7		
								60°	90°	120°	60°	90°	120°	60°	90°	120°
Sim.	Penalty SR ↑	0.95	0.96	0.56	0.85	Sim.	SR ↑	0.71	0.77	0.65	0.66	0.72	0.55	0.68	0.64	0.43
	Free-kick SR ↑	1.00	0.88	0.71	0.85		$E_{vel} \downarrow$	0.69	0.74	0.91	0.74	0.73	0.80	0.71	0.81	0.97
	—	—	—	—	—		$R_{bib} \uparrow$	0.73	0.74	0.64	0.76	0.75	0.70	0.76	0.72	0.60
Real	Penalty SR ↑	0.68(26/38)	0.57(28/49)	0.55(33/60)	—	Real	SR ↑	Easy: 0.79(22/28)			Medium: 0.68(19/28)			Hard: 0.61(17/28)		

TABLE I: Quantitative results for goal-directed shooting and directional dribbling in simulation and on the real robot. Shooting entries and real-robot dribbling entries report success rate (SR) on a 0–1 scale; simulation dribbling reports SR, velocity error (E_{vel}), and ball-in-band ratio (R_{bib}) over 900 repeated-S slalom trials. Arrows indicate the preferred direction, and bold marks the best result within each fixed- N simulation-dribbling group. [ZXX: For shooting, a success requires the ball to cross the goal line in one kick. For dribbling, a success requires the robot to remain upright while both robot and ball finish within 1 m of the target.]

Fig. 6. We record the success rate of the task. Experimental results in Table I demonstrating strong positional generalization capability. By analyzing the kinematic pattern during the shooting process in Fig. 6, we observe that the robot exhibits a human-like gait pattern.

Dribbling: The robot is required to complete a repeated-S slalom task. N is the number of poles, and α is the turning angle of instruction. A task is considered successful if the robot passes through all waypoints in the predefined order. To ensure consistency, we use a fixed instruction sequence for all experiments. The policy achieves an average success rate of 0.65 across 900 trials. Experimental results in Table I demonstrate that more than half of the task configurations obtain success rates above 0.65, showing robustness in long

trajectory and extreme angles. Visualization in Fig. 7 shows the dynamics of the policy.

Head-control comparison: We compare end-to-end learned head control with a heuristic that gazes at the predicted ball (Table II, last row). On static objects such as penalty kicks, the heuristic is comparable to learned active vision. On highly dynamic dribbling, the heuristic fails while the learned controller remains effective. The gap confirms that head motion must be learned jointly with locomotion and contact, rather than as a detached look-at-ball module, motivating active vision within the end-to-end policy.

Comparison with an external protocol: Dribble Master [27] is the closest published system to our setting. Both use full-size humanoids with 2-DoF active heads and

onboard sensing only; however, its policy consumes detector-derived ball coordinates, whereas DAVIS consumes a depth image. We evaluate DAVIS on its two core dribbling tasks with 50 trials per condition. At the 1 m tolerance, DAVIS achieves target-reaching and obstacle-avoidance success rates of 0.88 and 0.96, compared with the reported Dribble Master rates of 0.87 and 0.93.

Experiments show that DAVIS achieves high performance in all tasks. These results demonstrate that our framework are able to handle highly dynamic locomotion and fine-grained soccer manipulation.

B. Ablation Study

To verify the necessity of key structures and answer Q2, we conduct ablation studies on three tasks in simulation and record different metrics in Table II. All methods share the same training parameters and evaluation settings.

Ablation on curriculum. We remove the ramp-up and scaling curriculum. All reward terms are activated with full weights from the beginning. Without the progressive curriculum, the policy requires more training iterations to converge, and performance is significantly worse than baseline.

Ablation on auxiliary head. Our policy incorporates an object position prediction head. To validate effectiveness, we mask the predicted positions fed into the actor, while retaining the network architecture and input dimension. We disable the aux loss and stop backpropagation. As expected, without the explicit predictions, model struggles to localize objects to make contact and success drops to zero. To further isolate each switch of the full auxiliary head, we disaggregate this into four finer variants: w/o anneal keeps $\alpha = 0$ (always predicted auxiliary features; no GT $\rightarrow$ pred schedule); w/o geometry zeros the actor’s auxiliary input while keeping auxiliary losses; w/o supervision drops the goal/ball/visibility losses but still feeds predicted features; w/o vis. gate skips visibility $\times$ geometry masking. All components help success rate, and the full method is best.

Ablation on active vision. Active vision is another key design. [ZXX: Since its components all act through the same two head-joint control dimensions and cannot be ablated independently,] We fix the 2-DoF head joints, replace the original head-based rewards with body alignment rewards, and adjust the corresponding curriculum. After fixing the head joints, the ball frequently moves out of the FOV or enters blind spots, causing collapse. Another worth-noting observation is that the robot is unable to locate the gate. The performance drops nearly to zero, with a high possibility of falling.

C. Real-World Deployment

To answer Q3, we deploy our policies on the Noetix E1 humanoid robot equipped with a ZED2i depth camera on 2-DoF head. For the shooting experiments, we select 12 representative positions according to levels used in simulation. For dribbling experiments, we design three tasks including reaching target position, obstacle traversal, and

slalom dribbling. To better reflect real-world soccer scenarios, all experiments are conducted on a grass surface. Fig. 1 and Fig. 5 show that our policy could perform soccer task with high dynamics, precise actions. Quantitative results in Table I demonstrate that our policy maintains a high success rate in real deployment. Shooting records 26/38, 28/49, and 33/60 successes for the Easy, Medium, and Hard settings, respectively; a success requires the ball to cross the goal line in one kick. Dribbling records 22/28, 19/28, and 17/28 successes for target reaching, obstacle traversal, and slalom, respectively; a success requires the robot to remain upright while both robot and ball finish within 1 m of the target.

V. CONCLUSION

We presented a depth-only end-to-end active-vision framework DAVIS for humanoid soccer skills. The main idea is simple: keep the deployment interface minimal, but place enough structure inside training so that the policy can learn soccer-specific geometry, visibility, and whole-body contact behavior from depth and proprioception alone.

Within the shared framework, goal-directed shooting and directional dribbling serve as representative instances of one-shot striking and sustained ball control, retaining the observation-to-control interface while changing task-specific objects, commands, rewards, curricula, and motion-prior conditioning. Across simulation, ablations, and the complete Noetix E1 deployment pipeline, real-robot execution confirms that this depth-only interface transfers across contact regimes; together with additional skill demonstrations, these results support DAVIS as an extensible framework for humanoid soccer skills through task-specific training.

VI. LIMITATIONS AND FUTURE WORK

Two limitations remain. The first is perceptual: depth-only sensing keeps the interface compact but ties the policy to a single head-mounted camera, inheriting its occlusions, dropouts, and imperfectly simulated depth—a residual sim-to-real gap that limits real-world accuracy. The second sits upstream: the task-definition interface unifies shooting and dribbling, yet each skill needs hand-tuned rewards, curricula, and contact conditioning, leaving the pipeline a recipe, not an automatic one.

Closing the perceptual gap is largely engineering—sensor-noise modeling and real-world fine-tuning. Automating skill construction is the more open problem, where language-conditioned specification and automatic curricula could generate new behaviors on the same interface.

REFERENCES

- [1] N. Rudin, D. Hoeller, P. Reist, and M. Hutter, “Learning to walk in minutes using massively parallel deep reinforcement learning,” in *Proceedings of the 5th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 164. PMLR, 2022, pp. 91–100.
- [2] T. Miki, J. Lee, J. Hwangbo, L. Wellhausen, V. Koltun, and M. Hutter, “Learning robust perceptive locomotion for quadrupedal robots in the wild,” *Science Robotics*, vol. 7, no. 62, p. eabk2822, 2022.
- [3] Z. Zhuang, S. Yao, and H. Zhao, “Humanoid parkour learning,” 2024.

Method	Penalty kick			Free kick			Repeated-S slalom		
	SR $\uparrow$	E_{ang} $\downarrow$	$r_{contact}$ $\uparrow$	SR $\uparrow$	E_{ang} $\downarrow$	$r_{contact}$ $\uparrow$	SR $\uparrow$	E_{vel} $\downarrow$	r_{bib} $\uparrow$
DAVIS	0.84	23.29	99.75	0.80	30.18	99.27	0.70	0.79	0.74
(a) Ablation									
DAVIS-w/o-curriculum	0.62	27.40	32.53	0.55	34.42	87.67	0.51	0.77	0.69
DAVIS-auxiliary head-w/o-anneal	0.32	54.79	84.87	0.22	76.20	87.92	0.31	0.83	0.66
DAVIS-auxiliary head-w/o-geometry	0.20	53.53	87.33	0.15	69.47	78.13	0.34	0.91	0.54
DAVIS-auxiliary head-w/o-supervision	0.13	81.87	93.33	0.14	93.66	91.67	0.30	0.95	0.48
DAVIS-auxiliary head-w/o-visibility gate	0.34	51.97	91.33	0.29	68.44	92.08	0.27	1.06	0.64
DAVIS-w/o-auxiliary head	0	–	0	0	–	0	0	1.45	0.05
DAVIS-w/o-active vision	0	–	0	0.36	97.81	1.71	0	1.25	0.03
(b) Comparison									
DAVIS-Head controller	0.88	26.66	93.67	0.76	37.96	88.13	0.0	0.99	0.13

TABLE II: Simulation ablations and comparison. (a) Ablations. Penalty-kick and free-kick results use over 900 shooting trajectories, and repeated-S slalom uses over 100 dribbling trajectories. Metrics are success rate (SR), body-to-goal angular error (E_{ang}), successful ball-contact ratio ($r_{contact}$), velocity error (E_{vel}), and time in the control band (r_{bib}). The full method provides the best overall balance across tasks and metrics. (b) Comparison. Comparison between our end-to-end learned head control (Active vision) and a heuristic head controller (Heuristic) that gazes at the predicted ball. The heuristic performs comparably to active vision in static-object tasks (e.g., penalty kicks), while active vision shows a clear advantage in highly dynamic dribbling scenarios.

- [4] T. He, J. Gao, W. Xiao, Y. Zhang, Z. Wang, J. Wang, Z. Luo, G. He, N. Sobanbabu, C. Pan, Z. Yi, G. Qu, K. Kitani, J. Hodgins, L. Fan, Y. Zhu, C. Liu, and G. Shi, “Asap: Aligning simulation and real-world physics for learning agile humanoid whole-body skills,” 2025.
- [5] S. Zhu, Z. Zhuang, M. Zhao, K.-Y. Lee, and H. Zhao, “Hiking in the wild: A scalable perceptive parkour framework for humanoids,” 2026.
- [6] Z. Wu, X. Huang, L. Yang, Y. Zhang, K. Sreenath, X. Chen, P. Abbeel, R. Duan, A. Kanazawa, C. Sferrazza, G. Shi, and C. K. Liu, “Perceptive humanoid parkour: Chaining dynamic human skills via motion matching,” 2026.
- [7] H. Xiong, X. Xu, J. Wu, Y. Hou, J. Bohg, and S. Song, “Vision in action: Learning active perception from human demonstrations,” in *Proceedings of The 9th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 305. PMLR, 2025, pp. 5450–5463.
- [8] Y. Liu, S. Mu, X. Chao, Z. Li, Y. Mu, T. Chen, S. Li, C. Lyu, X.-P. Zhang, and W. Ding, “Avr: Active vision-driven robotic precision manipulation with viewpoint and focal length optimization,” 2025.
- [9] X. Cheng, J. Li, S. Yang, G. Yang, and X. Wang, “Open-television: Teleoperation with immersive active visual feedback,” 2024.
- [10] K. Nakagawa, Y. Ohmura, and Y. Kuniyoshi, “Imitation learning for active neck motion enabling robot manipulation beyond the field of view,” 2025.
- [11] Y. Fu, F. Xie, C. Xu, J. Xiong, H. Yuan, and Z. Lu, “Demohlm: From one demonstration to generalizable humanoid loco-manipulation,” 2025.
- [12] A. Kumar, Z. Fu, D. Pathak, and J. Malik, “Rma: Rapid motor adaptation for legged robots,” 2021.
- [13] H. Wang, H. Luo, W. Zhang, and H. Chen, “Cts: Concurrent teacher-student reinforcement learning for legged locomotion,” 2024.
- [14] H. Zhang, H. Yu, L. Zhao, A. Choi, Q. Bai, Y. Yang, and W. Xu, “Slim: Sim-to-real legged instructive manipulation via long-horizon visuomotor learning,” 2025.
- [15] H. Kitano, M. Asada, Y. Kuniyoshi, I. Noda, and E. Osawa, “Robocup: The robot world cup initiative,” in *Proceedings of the First International Conference on Autonomous Agents*, 1997, pp. 340–347.
- [16] R. Gerndt, D. Seifert, J. Baltes, S. Sadeghnejad, and S. Behnke, “Humanoid robots in soccer: Robots versus humans in robocup 2050,” *IEEE Robotics & Automation Magazine*, vol. 22, no. 3, pp. 147–154, 2015.
- [17] M. Beukman, B. Ingram, G. Nangue Tasse, B. Rosman, and P. Rancho, “Robocupgym: A challenging continuous control benchmark in robocup,” 2024.
- [18] S.-J. Yi, S. McGill, D. Hong, and D. D. Lee, “Hierarchical motion control for a team of humanoid soccer robots,” *International Journal of Advanced Robotic Systems*, vol. 13, no. 1, p. 32, 2016.
- [19] L. Leottau, C. Celemin, and J. Ruiz-del Solar, “Ball dribbling for humanoid biped robots: A reinforcement learning and fuzzy control approach,” in *RoboCup 2014: Robot World Cup XVIII*, ser. Lecture Notes in Computer Science, vol. 8992. Springer, 2015, pp. 549–561.
- [20] I. J. da Silva, D. H. Perico, T. P. D. Homem, and R. A. C. Bianchi, “Deep reinforcement learning for a humanoid robot soccer player,” *Journal of Intelligent & Robotic Systems*, vol. 102, no. 3, p. 69, 2021.
- [21] M. Abreu, L. P. Reis, and N. Lau, “Designing a skilled soccer team for robocup: Exploring skill-set-primitives through reinforcement learning,” *Neural Computing and Applications*, 2025.
- [22] Y. Ji, Z. Li, Y. Sun, X. B. Peng, S. Levine, G. Berseth, and K. Sreenath, “Hierarchical reinforcement learning for precise soccer shooting skills using a quadrupedal robot,” in *2022 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 2022, pp. 1479–1486.
- [23] X. Huang, Z. Li, Y. Xiang, Y. Ni, Y. Chi, Y. Li, L. Yang, X. B. Peng, and K. Sreenath, “Creating a dynamic quadrupedal robotic goalkeeper with reinforcement learning,” in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2023, pp. 2715–2722.
- [24] Y. Ji, G. B. Margolis, and P. Agrawal, “Dribblebot: Dynamic legged manipulation in the wild,” in *IEEE International Conference on Robotics and Automation (ICRA)*, 2023.
- [25] T. Haarnoja, B. Moran, G. Lever, S. H. Huang, D. Tirumala, J. Humpalik, M. Wulfmeier *et al.*, “Learning agile soccer skills for a bipedal robot with deep reinforcement learning,” *Science Robotics*, vol. 9, no. 89, p. eadi8022, 2024.
- [26] D. Tirumala, M. Wulfmeier, B. Moran, S. Huang, J. Humpalik, G. Lever, T. Haarnoja, L. Hasenclever, A. Byravan, N. Batchelor, N. Sreendra, K. Patel, M. Gwira, F. Nori, M. Riedmiller, and N. Heess, “Learning robot soccer from egocentric vision with deep reinforcement learning,” in *Proceedings of The 8th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 270, 2025, pp. 165–184.
- [27] Z. Wang, J. Zhou, and Q. Wu, “Dribble master: Learning agile humanoid dribbling through legged locomotion,” 2025.
- [28] J. Kong, X. Liu, Y. Lin, J. Han, S. Schwertfeger, C. Bai, and X. Li, “Learning soccer skills for humanoid robots: A progressive perception-action framework,” 2026.
- [29] Y. Chen, Z. Zhang, Z. Cao, Y. Chen, S. Fu, L. Yan, Y. Zhang, J. Liu, H. Li, and Y. Gao, “Hierkick: Hierarchical reinforcement learning for vision-guided soccer robot control,” 2026.
- [30] Y. Wang, C. Luo, P. Chen, J. Liu, W. Sun, T. Guo, K. Yang, B. Hu, Y. Zhang, and M. Zhao, “Learning vision-driven reactive soccer skills for humanoid robots,” 2025.
- [31] L. Sentis and O. Khatib, “A whole-body control framework for humanoids operating in human environments,” in *IEEE International Conference on Robotics and Automation (ICRA)*, 2006, pp. 2641–2648.
- [32] A. Herzog, N. Rotella, S. Mason, F. Grimmering, S. Schaal, and L. Righetti, “Momentum control with hierarchical inverse dynamics on a torque-controlled humanoid,” *Autonomous Robots*, vol. 40, no. 3, pp. 473–491, 2016.
- [33] T. He, W. Xiao, T. Lin, Z. Luo, Z. Xu, Z. Jiang, C. Liu, G. Shi, X. Wang, L. Fan, and Y. Zhu, “Hover: Versatile neural whole-body

- controller for humanoid robots,” *arXiv preprint arXiv:2410.21229*, 2024.
- [34] Y. Li, Y. Zhang, W. Xiao, C. Pan, H. Weng, G. He, T. He, and G. Shi, “Hold my beer: Learning gentle humanoid locomotion and end-effector stabilization control,” 2025.
 - [35] Q. Ben, F. Jia, J. Zeng, J. Dong, D. Lin, and J. Pang, “Homie: Humanoid loco-manipulation with isomorphic exoskeleton cockpit,” 2025.
 - [36] Y. Xue, W. Dong, M. Liu, W. Zhang, and J. Pang, “A unified and general humanoid whole-body controller for versatile locomotion,” 2025.
 - [37] W. Sun, L. Feng, B. Cao, Y. Liu, Y. Jin, and Z. Xie, “Ulc: A unified and fine-grained controller for humanoid loco-manipulation,” 2025.
 - [38] S. Bai, M. Li, X. Lv, J. Wang, X. Wang, F. Liao, C. Hou, L. Gu, W. Zhou, K. Wu, Z. Ding, Z. Xu, L. Sun, S. Zhang, Z. Che, J. Tang, and B. Chen, “Hex: Humanoid-aligned experts for cross-embodiment whole-body manipulation,” 2026.
 - [39] L. P. Kaelbling, M. L. Littman, and A. R. Cassandra, “Planning and acting in partially observable stochastic domains,” *Artificial Intelligence*, vol. 101, no. 1–2, pp. 99–134, 1998.
 - [40] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” 2017.
 - [41] X. B. Peng, P. Abbeel, S. Levine, and M. van de Panne, “Deepmimic: Example-guided deep reinforcement learning of physics-based character skills,” *ACM Transactions on Graphics*, vol. 37, no. 4, pp. 1–14, 2018.
 - [42] X. B. Peng, Z. Ma, P. Abbeel, S. Levine, and A. Kanazawa, “Amp: Adversarial motion priors for stylized physics-based character control,” *ACM Transactions on Graphics*, vol. 40, no. 4, pp. 1–20, 2021.
 - [43] A. Escontrela, X. B. Peng, W. Yu, T. Zhang, A. Iscen, K. Goldberg, and P. Abbeel, “Adversarial motion priors make good substitutes for complex reward functions,” 2022.
 - [44] J. Tobin, R. Fong, A. Ray, J. Schneider, W. Zaremba, and P. Abbeel, “Domain randomization for transferring deep neural networks from simulation to the real world,” in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2017, pp. 23–30.
 - [45] X. B. Peng, M. Andrychowicz, W. Zaremba, and P. Abbeel, “Sim-to-real transfer of robotic control with dynamics randomization,” in *IEEE International Conference on Robotics and Automation (ICRA)*, 2018, pp. 1–8.
 - [46] J. Long, Z. Wang, Q. Li, L. Cao, J. Gao, and J. Pang, “Hybrid internal model: Learning agile legged locomotion with simulated robot response,” in *The Twelfth International Conference on Learning Representations*, 2024.